

# Fake News Detection

## Introduction:

The world is filled with a lot of news generated everyday. News contains very crucial information, which has significant effect on public, government, economy of a country. In these modern days, fake news are very common, which lead to bad impact on social health. So, identifying fake news is very important. Manual checking is not suitable because of lots of manpower and cost.

This project is aimed to build a ML model which uses statistical analysis and modern machine learning approaches to predict the fake news using the text and title of the news.

## Dataset:

train.csv: A full training dataset with the following attributes:

id: unique id for a news article.

title: the title of a news article.

author: author of the news article.

text: the text of the article; could be incomplete.

label: a label that marks the article as potentially unreliable.

Where 1: unreliable and 0: reliable

## Objectives:

- Build a model which predicts the fake news using text and title of the news.
- To make a black list of words which are more likely to cause fake news.
- To determine whether there is difference in length of news (count of words) between real and fake news.

## Scope Of Study:

Fake news is very common in these days and it is very problematic. It will cause very serious situations and mislead people. And it is not easy to check whether the news is fake or not.

Machine learning approach to this will solve the problem easily. Without much cost, energy and time, ML models can be used to detect fake news more accurately.

## Exploratory Data Analysis

```
In [1]: #import libraries:

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
import re
from sklearn.ensemble import RandomForestRegressor, ExtraTreesRegressor
from sklearn.model_selection import train_test_split
from sklearn.model_selection import GridSearchCV
from sklearn.neighbors import KNeighborsRegressor
from sklearn.metrics import r2_score
from sklearn.linear_model import LogisticRegression
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.model_selection import train_test_split
from scipy.sparse import hstack
from sklearn.metrics import accuracy_score, confusion_matrix, log_loss, f1_score
import warnings
warnings.filterwarnings("ignore")
```

```
In [2]: data=pd.read_csv("D:\\collage project\\FAKE NEWS\\train.csv")
data
```

Out[2]:

	id	title	author	text	label
0	0	House Dem Aide: We Didn't Even See Comey's Let...	Darrell Lucas	House Dem Aide: We Didn't Even See Comey's Let...	1
1	1	FLYNN: Hillary Clinton, Big Woman on Campus - ...	Daniel J. Flynn	Ever get the feeling your life circles the rou...	0
2	2	Why the Truth Might Get You Fired	Consortiumnews.com	Why the Truth Might Get You Fired October 29, ...	1
3	3	15 Civilians Killed In Single US Airstrike Hav...	Jessica Purkiss	Videos 15 Civilians Killed In Single US Aistr...	1
4	4	Iranian woman jailed for fictional unpublished...	Howard Portnoy	Print \nAn Iranian woman has been sentenced to...	1
...	...	...	...	...	...
20795	20795	Rapper T.I.: Trump a 'Poster Child For White S...	Jerome Hudson	Rapper T. I. unloaded on black celebrities who...	0
20796	20796	N.F.L. Playoffs: Schedule, Matchups and Odds -...	Benjamin Hoffman	When the Green Bay Packers lost to the Washing...	0
20797	20797	Macy's Is Said to Receive Takeover Approach by...	Michael J. de la Merced and Rachel Abrams	The Macy's of today grew from the union of sev...	0
20798	20798	NATO, Russia To Hold Parallel Exercises In Bal...	Alex Ansary	NATO, Russia To Hold Parallel Exercises In Bal...	1
20799	20799	What Keeps the F-35 Alive	David Swanson	David Swanson is an author, activist, journa...	1

20800 rows × 5 columns

```
In [3]: data.head()
```

Out[3]:

	id	title	author	text	label
0	0	House Dem Aide: We Didn't Even See Comey's Let...	Darrell Lucas	House Dem Aide: We Didn't Even See Comey's Let...	1
1	1	FLYNN: Hillary Clinton, Big Woman on Campus - ...	Daniel J. Flynn	Ever get the feeling your life circles the rou...	0

	id	title	author	text	label
2	2	Why the Truth Might Get You Fired	Consortiumnews.com	Why the Truth Might Get You Fired October 29, ...	1
3	3	15 Civilians Killed In Single US Airstrike Hav...	Jessica Purkiss	Videos 15 Civilians Killed In Single US Aistr...	1
4	4	Iranian woman jailed for fictional unpublished...	Howard Portnoy	Print \nAn Iranian woman has been sentenced to...	1

In [4]:

```
# Shape of the Data:  
data.shape
```

Out[4]: (20800, 5)

This Dataset contains 5 variables and 20800 observations.

In [5]:

```
# information About the Dataset:  
  
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 20800 entries, 0 to 20799  
Data columns (total 5 columns):  
#   Column  Non-Null Count  Dtype  
---  -  
0    id      20800 non-null    int64  
1   title   20242 non-null    object  
2  author  18843 non-null    object  
3   text   20761 non-null    object  
4  label   20800 non-null    int64  
dtypes: int64(2), object(3)  
memory usage: 812.6+ KB
```

In [6]:

```
# Data types:  
data.dtypes
```

Out[6]: id int64  
title object  
author object  
text object  
label int64  
dtype: object

In [7]:

```
data.describe()
```

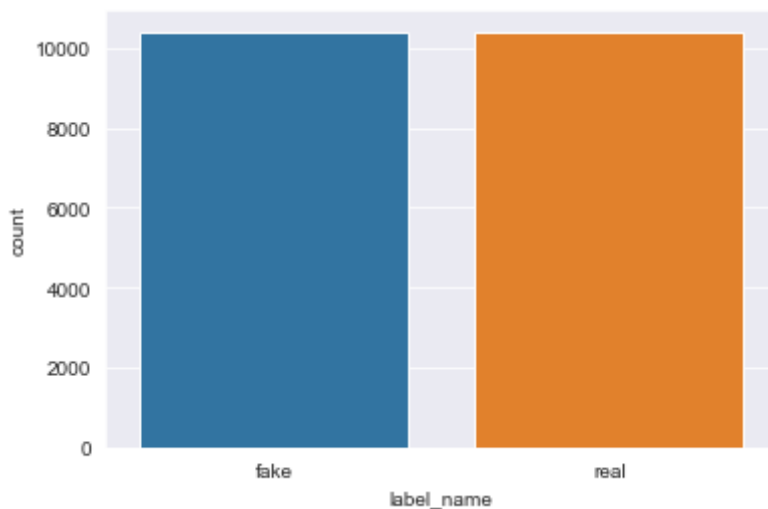
Out[7]:

	id	label
count	20800.000000	20800.000000
mean	10399.500000	0.500625
std	6004.587135	0.500012
min	0.000000	0.000000
25%	5199.750000	0.000000
50%	10399.500000	1.000000
75%	15599.250000	1.000000
max	20799.000000	1.000000

## Check the imbalance

```
In [8]: # check whether our data is balanced :
data["label_name"] = np.where(data["label"] == 1, "fake", "real")
sns.set_style('darkgrid')
sns.countplot(data['label_name'])
```

```
Out[8]: <AxesSubplot:xlabel='label_name', ylabel='count'>
```



Here 0 represent Fake News and 1 represent Real news

We can conclude from the plot that the data is balanced.

This is important because, without knowing the data imbalance, we cannot consider the appropriate metric and model.

## Choosing the performance metric:

### Accuracy:

Accuracy can be used since the data is balanced. And it is simple to understand and interpret.

## Data Cleaning:

```
In [9]: data.drop("label_name", axis=1, inplace=True)
```

```
In [10]: # Find Null Values:
null = data.isnull().sum()
null
```

```
Out[10]: id          0
title       558
author     1957
text        39
label       0
dtype: int64
```

Title and Author and Text columns having lot of null values so fill null values by imputation method.

```
In [11]: # Columns Name:
data.columns
```

```
Out[11]: Index(['id', 'title', 'author', 'text', 'label'], dtype='object')
```

# Data Preprocessing:

```
In [12]: #calculating total null values:

total_rows =data.shape[0]
total_null =null.sum()

# percent of data that is null:

print("Percentage of null value instances present is : ",round((total_null/total_row

Percentage of null value instances present is : 12.279 %
```

```
In [13]: # Impute missing values in dish liked with "unknown", so that "unknown" also conside
data['author'] = data['author'].fillna("unknown")
data['author']
```

```
Out[13]: 0          Darrell Lucas
1          Daniel J. Flynn
2      Consortiumnews.com
3          Jessica Purkiss
4          Howard Portnoy
...
20795          Jerome Hudson
20796      Benjamin Hoffman
20797      Michael J. de la Merced and Rachel Abrams
20798          Alex Ansary
20799      David Swanson
Name: author, Length: 20800, dtype: object
```

```
In [14]: #calculating total null values:

total_rows = data.shape[0]
total_null = null.sum()

# percent of data that is null:

print("Percentage of null value instances present is : ",round((total_null/total_row

Percentage of null value instances present is : 12.279 %
```

```
In [15]: # Remove all other Null Values:
data.dropna(inplace=True)
```

```
In [16]: data.isnull().sum()
```

```
Out[16]: id          0
title          0
author         0
text           0
label          0
dtype: int64
```

So there is no null values now.

## Text Cleaning:

```
In [17]: # https://stackoverflow.com/a/47091490/4084039

def decontracted(phrase):
    """This function will converts shorthend form to full form of english words"""
    # specific
```

```

phrase = re.sub(r"won't", "will not", phrase)
phrase = re.sub(r"can't", "can not", phrase)

# general
phrase = re.sub(r"n't", " not", phrase)
phrase = re.sub(r"'re", " are", phrase)
phrase = re.sub(r"'s", " is", phrase)
phrase = re.sub(r"'d", " would", phrase)
phrase = re.sub(r"'ll", " will", phrase)
phrase = re.sub(r"'t", " not", phrase)
phrase = re.sub(r"'ve", " have", phrase)
phrase = re.sub(r"'m", " am", phrase)
return phrase

```

```

In [18]: from tqdm import tqdm
def preprocess_text(text_data):
    """This function will clean the text data"""
    preprocessed_text = []
    # tqdm is for printing the status bar
    for sentence in tqdm(text_data):
        sent = decontracted(sentence)

        # removing \r
        sent = sent.replace('\r', ' ')

        #remove NaN text
        sent = sent.replace('NaN', ' ')

        #
        sent = sent.replace('\n', ' ')

        #
        sent = sent.replace('\\"', ' ')

        #substituting number with space
        sent = re.sub(r'[0-9]+', ' ',sent)

        sent = re.sub(r'[A-Za-z0-9]+', ' ', sent)

        sent = ' '.join([w for w in sent.split() if len(w)>=3])
        sent = ' '.join(e for e in sent.split() if e.lower() not in stopwords)
        preprocessed_text.append(sent.lower().strip())
        # https://gist.github.com/sebleier/554280

    return preprocessed_text

```

```

In [19]: def categorical_data(cat):

    """This function cleanse categorical text data"""

    preprocessed_cat = []

    for point in cat:

        lst_words = point.split()

        for i in range(len(lst_words)):
            splitted =lst_words[i].split()

            if len(splitted)>1:
                lst_words[i] = "_".join(splitted)

```

```
point = "".join(1st_words).lower()

preprocessed_cat.append(point)

return preprocessed_cat
```

```
In [20]: # https://gist.github.com/sebleier/554280
# we are removing the words from the stop words list: 'no', 'nor', 'not' because the

stopwords= ['i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves', 'you', "you",
            "you'll", "you'd", 'your', 'yours', 'yourself', 'yourselves', 'he', 'him',
            'she', "she's", 'her', 'hers', 'herself', 'it', "it's", 'its', 'itself',
            'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'this', 'that',
            'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'have', 'has',
            'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or', 'because', 'at',
            'by', 'for', 'with', 'about', 'against', 'between', 'into', 'throu',
            'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out', 'on', 'off',
            'then', 'once', 'here', 'there', 'when', 'where', 'why', 'how', 'all', 't',
            'most', 'other', 'some', 'such', 'only', 'own', 'same', 'so', 'than', 't',
            's', 't', 'can', 'will', 'just', 'don', "don't", 'should', "should've",
            've', 'y', 'ain', 'aren', "aren't", 'couldn', "couldn't", 'didn', "didn't",
            "hadn't", 'hasn', "hasn't", 'haven', "haven't", 'isn', "isn't", 'ma', 'm',
            "mustn't", 'needn', "needn't", 'shan', "shan't", 'shouldn', "shouldn't",
            'won', "won't", 'wouldn', "wouldn't", "NaN",]
```

```
In [21]: data['text'] = preprocess_text(data['text'].values)
```

```
100%|██████████████████████████████████████████████████████████████████████████| 20  
203/20203 [01:50<00:00, 183.41it/s]
```

```
In [22]: data['title'] = preprocess_text(data['title'].values)
```

```
100%|██████████████████████████████████████████████████████████████████████████| 202  
03/2023 [00:02<00:00, 7716.12it/s]
```

```
In [23]: data['author'] = categorical_data(data['author'].values)
```

```
In [24]: data['text'][0]
```

Out[24]: 'house dem aide even see comej letter jason chaffetz tweeted darrell lucus october s  
ubscribe jason chaffetz stump american fork utah image courtesy michael jolley avail  
able creative commons license apologies keith olbermann doubt worst person world wee  
k fbi director james comej according house democratic aide looks like also know seco  
nd worst person well turns comej sent infamous letter announcing fbi looking emails  
may related hillary clinton email server ranking democrats relevant committees hear  
comej found via tweet one republican committee chairmen know comej notified republic  
an chairmen democratic ranking members house intelligence judiciary oversight commit  
tees agency reviewing emails recently discovered order see contained classified info  
rmation not long letter went oversight committee chairman jason chaffetz set politic  
al world ablaze tweet fbi dir informed fbi learned existence emails appear pertinent  
investigation case reopened jason chaffetz jasoninthehouse october course know not c  
ase comej actually saying reviewing emails light unrelated case know anthony weiner  
sexting teenager apparently little things facts matter chaffetz utah republican alre  
ady vowed initiate raft investigations hillary wins least two years worth possibly e  
ntire term worth apparently chaffetz thought fbi already work resulting tweet briefl  
y roiled nation cooler heads realized dud according senior house democratic aide mis  
reading letter may least chaffetz sins aide told shareblue boss democrats even know  
comej letter time found checked twitter democratic ranking members relevant committe  
es receive comej letter republican chairmen fact democratic ranking members receive  
chairman oversight government reform committee jason chaffetz tweeted made public le  
t see got right fbi director tells chaffetz gop committee chairmen major development  
potentially politically explosive investigation neither chaffetz nor colleagues cour  
tesy let democratic counterparts know instead according aide made find twitter alrea  
dy talk daily kos comej provided advance notice letter chaffetz republicans giving t  
ime turn spin machine may make good theater nothing far even suggests case nothing f

ar suggests comey anything grossly incompetent tone deaf suggest however chaffetz acting way makes dan burton darrell issa look like models responsibility bipartisanship even decency notify ranking member elijah cummings something explosive trample basic standards fairness know granted not likely chaffetz answer sits ridiculously republican district anchored provo orem cook partisan voting index gave mitt romney punishing percent vote moreover republican house leadership given full support chaffetz planned fishing expedition mean turn hot lights textbook example house become republican control also second worst person world darrell lucus darrell something graduate university north carolina considers journalist old school attempt turn member religious right college succeeded turning religious right worst nightmare charismatic christian unapologetic liberal desire stand scared silence increased survived abusive three year marriage may know daily kos christian dem follow twitter darrelllucuss connect facebook click buy darrell mello yello connect'

In [25]: `data['title'][0]`

Out[25]: 'house dem aide even see comey letter jason chaffetz tweeted'

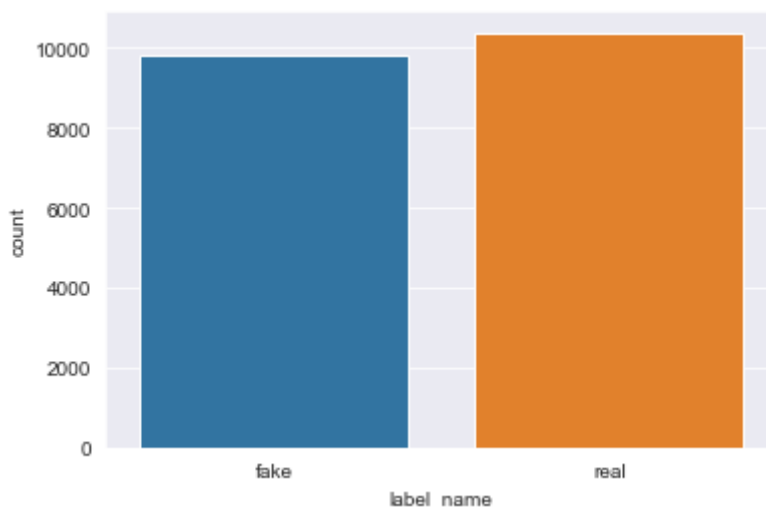
In [26]: `data['author'][0]`

Out[26]: 'darrelllucuss'

In [27]: 

```
# check whether our data is balanced :
data["label_name"] = np.where(data["label"] == 1, "fake", "real")
sns.set_style('darkgrid')
sns.countplot(data['label_name'])
```

Out[27]: `<AxesSubplot:xlabel='label_name', ylabel='count'>`



In [28]: `data['label_name'].value_counts()`

Out[28]: 

```
real    10387
fake     9816
Name: label_name, dtype: int64
```

## Model Splitting:

In [29]: 

```
#train test splitting:
y = data['label']
x = data.drop(["label", "id"], axis=1)
```

In [30]: `x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=.2, random_state=353)`

In [31]: `x_train.shape`

Out[31]: (16162, 4)



```
In [32]: x_test.shape
```

```
Out[32]: (4041, 4)
```

```
In [33]: y_train.shape
```

```
Out[33]: (16162,)
```

```
In [34]: y_test.shape
```

```
Out[34]: (4041,)
```

## BOW vectorizer:

```
In [35]: vec = CountVectorizer()

#fitting countvectorizer using only train data:

vec.fit(x_train["text"].values)

#transforming to vector representation for train,test data:

x_train_text = vec.transform(x_train["text"].values)
x_test_text = vec.transform(x_test["text"].values)

print(x_train_text.shape)
print(x_test_text.shape)
```

```
(16162, 132829)
(4041, 132829)
```

```
In [36]: vec = CountVectorizer()

#fitting countvectorizer using only train data:

vec.fit(x_train["title"].values)

#transforming to vector representation for train,test data:

x_train_title = vec.transform(x_train["title"].values)
x_test_title = vec.transform(x_test["title"].values)

print(x_train_title.shape)
print(x_test_title.shape)
```

```
(16162, 18539)
(4041, 18539)
```

```
In [37]: # Use one hot encoding for this

vec = CountVectorizer()

#fitting countvectorizer using only train data:

vec.fit(x_train["author"].values)

#transforming to vector representation for train,test data:

x_train_author = vec.transform(x_train["author"].values)
x_test_author = vec.transform(x_test["author"].values)
```

```
print(x_train_author.shape)
print(x_test_author.shape)
```

```
(16162, 3668)
(4041, 3668)
```

In [38]: `from scipy.sparse import hstack`

In [39]: `#Concatenating all features
x_tr=hstack((x_train_text,x_train_title,x_train_author)).tocsr()
x_te=hstack((x_test_text,x_test_title,x_test_author)).tocsr()

print("FINAL DATA MATRIX SHAPE IS .....")
print(x_tr.shape,y_train.shape)
print(x_te.shape,y_test.shape)
print("*"*100)`

```
FINAL DATA MATRIX SHAPE IS .....
(16162, 155036) (16162,)
(4041, 155036) (4041,)
*****
*****
```

## Data Analysis:

In [133... `# determine whether title contain more number of real or fake news.
fake_news_count = data[data.label == 1]['title'].value_counts()
real_news_count = data[data.label == 0]['title'].value_counts()`

In [134... `fake_count = pd.DataFrame({
 'title':fake_news_count.index,
 'Fake':fake_news_count.values
})
real_count = pd.DataFrame({
 'title':real_news_count.index,
 'Real':real_news_count.values
})`

In [135... `rf_count = pd.merge(real_count, fake_count, on='title', how='outer').fillna(0)
rf_count['Real'] = rf_count['Real'].astype(int)
rf_count['Fake'] = rf_count['Fake'].astype(int)
rf_count`

Out[135...

	title	Real	Fake
0	cook week new york times	4	0
1	great stories nothing politics new york times	4	0
2	right left partisan writing miss new york times	2	0
3	trump new york times	2	0
4	cook weekend new york times	2	0
...	...	...	...
19346	comments week comes trouble	0	1
19347	comment holocaust jewish population numbers je...	0	1
19348	hillary clinton private speech mentioned pales...	0	1
19349	lawsuits states accuse trump gop voter intimid...	0	1

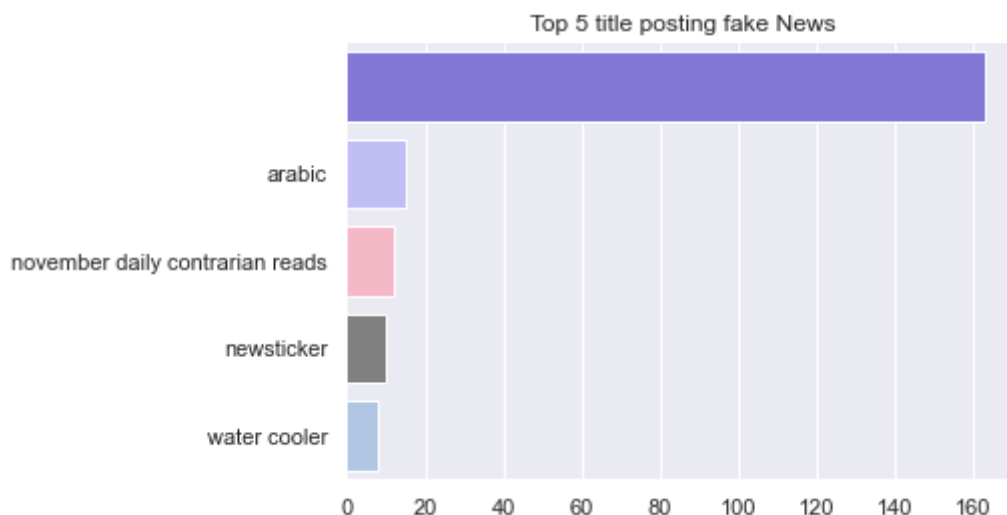
	title	Real	Fake
19350	took less hours new email investigation story ...	0	1

19351 rows × 3 columns

In [145...

```
#Cheack top five Title contains Fake news:

sns.barplot(y=fake_news_count[:5].index, x=fake_news_count[:5].values,
            palette=['#7868e6', '#b8b5ff', '#ffaec0', 'grey', '#a7c5eb'])
sns.despine(bottom=True, left=True)
plt.title('Top 5 title posting fake News');
```



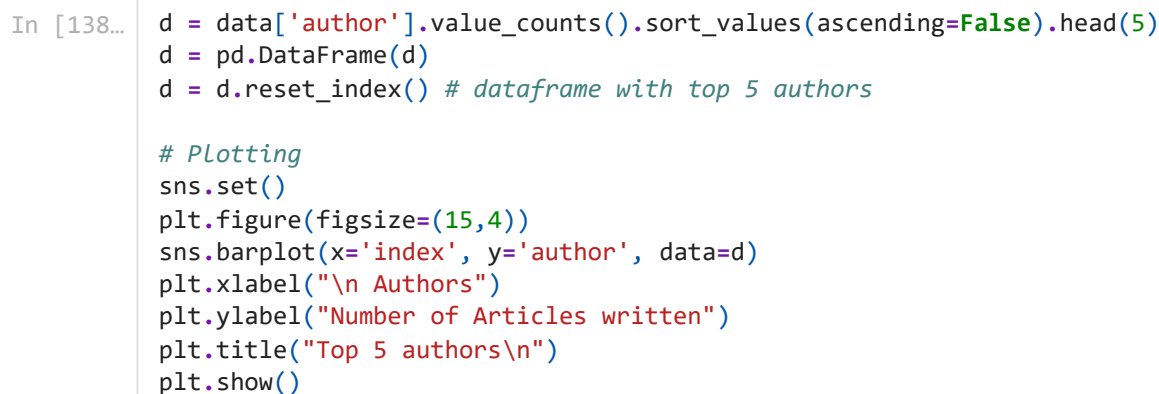
Here "Cook week new york times" and "Great stories nothing politics new york times" titles having more fake news frequency

In [137...

```
from wordcloud import WordCloud, STOPWORDS, ImageColorGenerator
stopwords = set(STOPWORDS)

wordcloud = WordCloud(width = 800, height = 800,
                      background_color = 'black',
                      stopwords = stopwords,
                      min_font_size = 10).generate(str(data['text']))

# plot the WordCloud image
plt.figure(figsize = (8,8), facecolor = 'blue')
plt.imshow(wordcloud)
plt.axis("off")
plt.tight_layout(pad = 0) ;
```



Here Pam key and admin authors are high.

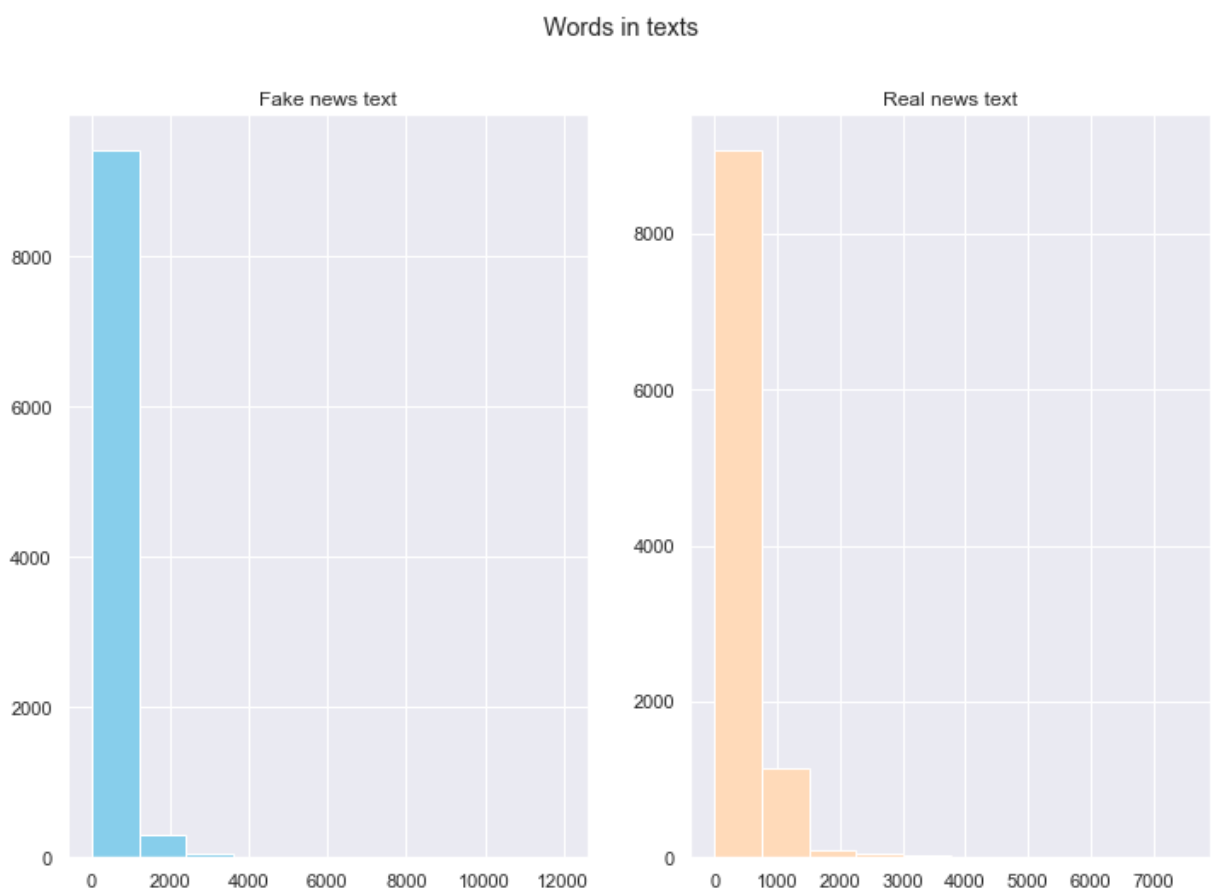
There are lot of news which do not contain author names i.e. unknown

## Analysis of count of words in text

```
In [139... # Number of words in each text.

fig,(ax1,ax2)=plt.subplots(1,2,figsize=(12,8))
text_len=data[data['label']==1]['text'].str.split().map(lambda x: len(x))
ax1.hist(text_len,color='SkyBlue')
ax1.set_title('Fake news text')

text_len=data[data['label']==0]['text'].str.split().map(lambda x: len(x))
ax2.hist(text_len,color='PeachPuff')
ax2.set_title('Real news text')
fig.suptitle('Words in texts')
plt.show()
```



1000 words are most common in real news category while around 500 words are most common in fake news category.

By just observing this, we can say that real news contains more words than fake news. But this is not a statistical conclusion. We can carry out 2 sample t-test to compare this.

```
In [140... data["length"] = [len(text) for text in data["text"]]

real_news_sample = data[data["label"]==0]["length"]
fake_news_sample = data[data["label"]==1]["length"]
```

```
In [141... real_news_sample.mean()
```



15/32



```

{'C': 0.1},
{'C': 1},
{'C': 10}],
'split0_test_score': array([0.86545005, 0.93751933, 0.96814105, 0.97680173, 0.97680
173,
    0.97680173]),
'split1_test_score': array([0.8731828 , 0.94401485, 0.97092484, 0.97742035, 0.97711
104,
    0.97649242]),
'split2_test_score': array([0.86881188, 0.93997525, 0.96751238, 0.97431931, 0.97370
05 ,
    0.97215347]),
'split3_test_score': array([0.87066832, 0.94306931, 0.96905941, 0.97462871, 0.97524
752,
    0.97462871]),
'split4_test_score': array([0.87066832, 0.93842822, 0.96844059, 0.97215347, 0.97153
465,
    0.97215347]),
'mean_test_score': array([0.86975627, 0.94060139, 0.96881565, 0.97506471, 0.9748790
9,
    0.97444596]),
'std_test_score': array([0.00256282, 0.00254386, 0.00116639, 0.00188613, 0.0020682
5,
    0.00201406]),
'rank_test_score': array([6, 5, 4, 1, 2, 3])}

```

```

In [56]: plt.plot(hyperparameter["C"], search.cv_results_['mean_test_score'])
plt.xscale("log")
plt.xlabel("Hyper parameter (C)")
plt.ylabel("Accuracy score")
plt.title("Hyper parameter v/s Accuracy score")
plt.show()

```



For  $c=0.1$  validation accuracy is high. so we can consider 0.1 is the best hyperparameter for our model.

```

In [57]: search.best_params_

```

```

Out[57]: {'C': 0.1}

```

## Model with best hyperparameter

```

In [58]: model = LogisticRegression(C = search.best_params_["C"])
model.fit(x_tr,y_train)

```

```

Out[58]:

```



▼ LogisticRegression

LogisticRegression(C=0.1)

```
In [59]: coefs_df = pd.DataFrame()
coefs_df["coefs"] = model.coef_[0]
coefs_df["abs_coefs"] = abs(model.coef_[0])
coefs_df.sort_values(by="abs_coefs", inplace=True, ascending=False)
```

```
In [60]: top_coefs = coefs_df.head(50)
top_coefs
```


Out[60]:

	coefs	abs_coefs
134846	-2.850448	2.850448
151284	-1.643217	1.643217
149518	-1.598940	1.598940
154865	1.429272	1.429272
151861	-1.096006	1.096006
14985	-1.039914	1.039914
143895	-1.010139	1.010139
42900	-1.002556	1.002556
82664	0.962555	0.962555
121452	-0.834175	0.834175
5131	0.795343	0.795343
151860	-0.772029	0.772029
81636	0.753945	0.753945
35637	0.639349	0.639349
80970	0.592978	0.592978
151862	-0.576345	0.576345
152002	-0.571661	0.571661
106180	0.556984	0.556984
110106	0.515807	0.515807
52671	0.511006	0.511006
22402	0.491886	0.491886
125800	0.482916	0.482916
153216	-0.479944	0.479944
71281	-0.457014	0.457014
105024	0.442244	0.442244
113843	-0.430602	0.430602
103119	-0.413177	0.413177
77387	-0.404610	0.404610

	coefs	abs_coefs
<b>82652</b>	0.399795	0.399795
<b>83153</b>	0.383349	0.383349
<b>152096</b>	0.380029	0.380029
<b>151429</b>	0.362534	0.362534
<b>91801</b>	-0.361924	0.361924
<b>109128</b>	0.348876	0.348876
<b>92301</b>	0.348701	0.348701
<b>59893</b>	-0.338336	0.338336
<b>59104</b>	-0.335731	0.335731
<b>1744</b>	-0.333319	0.333319
<b>92168</b>	0.333046	0.333046
<b>76303</b>	-0.319583	0.319583
<b>21378</b>	0.318092	0.318092
<b>151863</b>	-0.317486	0.317486
<b>22621</b>	0.316044	0.316044
<b>106196</b>	0.311680	0.311680
<b>47093</b>	0.307559	0.307559
<b>102232</b>	-0.303310	0.303310
<b>117989</b>	-0.297028	0.297028
<b>69979</b>	-0.293608	0.293608
<b>5882</b>	-0.292243	0.292243
<b>136016</b>	0.287326	0.287326

```
In [61]: top_coefs["feature index"] = top_coefs.index
```

```
In [62]: top_coefs[["coefs"]].plot(kind="bar",figsize=(15, 6))
plt.title("Feature importance plot")
plt.xlabel("Feature index")
plt.ylabel("Co-efficient value")
plt.show()
```



In [ ]:

In [ ]:

In [ ]:

## Model Evaluation

In [ ]:

```
In [63]: y_test_pred = model.predict(x_te)
print("TEST PREDICTION: \n",y_test_pred[:3])
#predicting probability
y_test_pred_proba = model.predict_proba(x_te)
print("TEST PREDICTION PROBABILITY: \n",y_test_pred_proba[:3])

train_pred = model.predict(x_tr)
#predicting probability
train_pred_proba = model.predict_proba(x_tr)

accuracy = accuracy_score(y_test,y_test_pred)
print("Test log loss of the model is : ", np.round(accuracy*100,4), "%")
train_accuracy = accuracy_score(y_train,train_pred)
print("Train log loss of the model is : ", np.round(train_accuracy*100,4), "%")

print("="*50)

test_logloss = log_loss(y_test,y_test_pred_proba)
print("Test log loss of the model is : ", np.round(test_logloss,4))
train_logloss = log_loss(y_train,train_pred_proba)
print("Train log loss of the model is : ", np.round(train_logloss,4))

print("="*50)

test_f1 = f1_score(y_test,y_test_pred)
print("Test f1 score of the model is : ", np.round(test_f1*100,4))
train_f1 = f1_score(y_train,train_pred)
print("Train f1 score of the model is : ", np.round(train_f1*100,4))

cm1 = confusion_matrix(y_test,y_test_pred)
```

```
plt.figure(figsize=(7,5))
sns.heatmap(cm1, annot=True,fmt="d",cmap='Blues')
plt.xlabel("predicted")
plt.ylabel("actual")
plt.title("Confusion matrix")
plt.show()
```

TEST PREDICTION:

[0 1 1]

TEST PREDICTION PROBABILITY:

[[0.89573426 0.10426574]

[0.01157584 0.98842416]

[0.00109493 0.99890507]]

Test log loss of the model is : 97.8966 %

Train log loss of the model is : 99.9567 %

=====


Test log loss of the model is : 0.0746

Train log loss of the model is : 0.015

=====

Test f1 score of the model is : 97.8541

Train f1 score of the model is : 99.9553



In [ ]:

- 2018 data points which belongs to fake news are classified correctly.
- 46 data points which belongs to real news category are misclassified as fake news.
- 39 data points which belongs to fake news category are misclassified as real news.
- 1938 data points which belongs to real news are classified correctly.

## KNN

```
In [64]: #https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LogisticRegr
from sklearn.neighbors import KNeighborsClassifier

#https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.Randomize
from sklearn.model_selection import RandomizedSearchCV
```

```
In [65]: model = KNeighborsClassifier()
```

```
In [66]: hyperparameter = {"n_neighbors" : [3,5,7,9,11,13,15,19,21]}
```

```
search = RandomizedSearchCV(model, hyperparameter, scoring = "accuracy", random_stat
search.fit(x_tr,y_train)
```


Out[66]:

```
RandomizedSearchCV
estimator: KNeighborsClassifier
KNeighborsClassifier
```

In [67]: search.cv\_results\_

```
Out[67]: {'mean_fit_time': array([0.05986347, 0.06069417, 0.06243091, 0.07088747, 0.06324778,
0.06376486, 0.06740904, 0.0750968 , 0.06753979]),
'std_fit_time': array([0.01237112, 0.00719888, 0.01192099, 0.00530294, 0.00661318,
0.01034121, 0.00897797, 0.00966789, 0.00558381]),
'mean_score_time': array([8.69746289, 9.42936983, 9.59136209, 9.806181 , 9.5899579
5,
9.21916165, 9.58306046, 9.5725194 , 9.70794082]),
'std_score_time': array([0.06309577, 0.18880388, 0.38274733, 0.5693801 , 0.1479986
9,
0.18249445, 0.44665977, 0.26740583, 0.07985271]),
'param_n_neighbors': masked_array(data=[3, 5, 7, 9, 11, 13, 15, 19, 21],
mask=[False, False, False, False, False, False, False, False, False],
fill_value='?',
dtype=object),
'params': [{ 'n_neighbors': 3},
{ 'n_neighbors': 5},
{ 'n_neighbors': 7},
{ 'n_neighbors': 9},
{ 'n_neighbors': 11},
{ 'n_neighbors': 13},
{ 'n_neighbors': 15},
{ 'n_neighbors': 19},
{ 'n_neighbors': 21}],
'split0_test_score': array([0.73368388, 0.7129601 , 0.70337148, 0.69873183, 0.69347
355,
0.68543149, 0.68079183, 0.67336839, 0.6705846 ]),
'split1_test_score': array([0.7370863 , 0.72038354, 0.71481596, 0.70832045, 0.70306
217,
0.69347355, 0.69223631, 0.68481287, 0.67924528]),
'split2_test_score': array([0.73019802, 0.72153465, 0.71441832, 0.70544554, 0.70266
089,
0.69740099, 0.69337871, 0.68873762, 0.68626238]),
'split3_test_score': array([0.73855198, 0.73174505, 0.72029703, 0.70946782, 0.70018
564,
0.69461634, 0.6927599 , 0.68347772, 0.67790842]),
'split4_test_score': array([0.75 , 0.73607673, 0.72957921, 0.72524752, 0.71875
,
0.7082302 , 0.70544554, 0.69863861, 0.69523515]),
'mean_test_score': array([0.73790404, 0.72454002, 0.7164964 , 0.70944263, 0.7036264
5,
0.69583051, 0.69292246, 0.68580704, 0.68184716]),
'std_test_score': array([0.00670214, 0.00831157, 0.00854413, 0.00873788, 0.0083053
3,
0.00736733, 0.00780622, 0.00817767, 0.008341 ]),
'rank_test_score': array([1, 2, 3, 4, 5, 6, 7, 8, 9])}
```

```
In [68]: plt.plot(hyperparameter["n_neighbors"], search.cv_results_['mean_test_score'])
plt.xlabel("Hyper parameter (n_neighbors)")
plt.ylabel("Accuracy score")
plt.title("Hyper parameter v/s Accuracy score")
plt.show()
```



In [69]: `search.best_estimator_`

Out[69]: `KNeighborsClassifier`  
`KNeighborsClassifier(n_neighbors=3)`

In [70]: `search.best_params_`

Out[70]: `{'n_neighbors': 3}`

In [71]: `model = KNeighborsClassifier(n_neighbors = search.best_params_["n_neighbors"])`  
`model.fit(x_tr,y_train)`

Out[71]: `KNeighborsClassifier`  
`KNeighborsClassifier(n_neighbors=3)`

In [ ]:

In [72]: `y_test_pred = model.predict(x_te)`

In [ ]:


In [73]: `from sklearn.metrics import accuracy_score, confusion_matrix`

In [74]: `accuracy = accuracy_score(y_test,y_test_pred)`  
`print("Test accuracy of the model is : ", np.round(accuracy*100,4), "%" )`  
`train_pred = model.predict(x_tr)`  
`train_accuracy = accuracy_score(y_train,train_pred)`  
`print("Train accuracy of the model is : ", np.round(train_accuracy*100,4), "%")`  
Test accuracy of the model is : 75.6248 %  
Train accuracy of the model is : 82.5393 %

In [75]: `cm1 = confusion_matrix(y_test,y_test_pred)`

In [76]: `plt.figure(figsize=(7,5))`  
`sns.heatmap(cm1, annot=True,fmt="d",cmap='Blues')`  
`plt.xlabel("predicted")`

```
plt.ylabel("actual")
plt.title("Confusion matrix")
plt.show()
```



```
In [77]: y_test_pred = model.predict(x_te)
print("TEST PREDICTION: \n",y_test_pred[:3])
#predicting probability
y_test_pred_proba = model.predict_proba(x_te)
print("TEST PREDICTION PROBABILITY: \n",y_test_pred_proba[:3])

train_pred = model.predict(x_tr)
#predicting probability
train_pred_proba = model.predict_proba(x_tr)

accuracy = accuracy_score(y_test,y_test_pred)
print("Test log loss of the model is : ", np.round(accuracy*100,4), "%")
train_accuracy = accuracy_score(y_train,train_pred)
print("Train log loss of the model is : ", np.round(train_accuracy*100,4), "%")

print("="*50)

test_logloss = log_loss(y_test,y_test_pred_proba)
print("Test log loss of the model is : ", np.round(test_logloss,4))
train_logloss = log_loss(y_train,train_pred_proba)
print("Train log loss of the model is : ", np.round(train_logloss,4))

print("="*50)

test_f1 = f1_score(y_test,y_test_pred)
print("Test f1 score of the model is : ", np.round(test_f1*100,4))
train_f1 = f1_score(y_train,train_pred)
print("Train f1 score of the model is : ", np.round(train_f1*100,4))

cm1 = confusion_matrix(y_test,y_test_pred)
plt.figure(figsize=(7,5))
sns.heatmap(cm1, annot=True,fmt="d",cmap='Blues')
plt.xlabel("predicted")
plt.ylabel("actual")
```

```
plt.title("Confusion matrix")
plt.show()
```

TEST PREDICTION:

```
[0 1 1]
```

TEST PREDICTION PROBABILITY:

```
[[1. 0.]
```

```
[0. 1.]
```

```
[0. 1.]]
```

Test log loss of the model is : 75.6248 %


Train log loss of the model is : 82.5393 %

Test log loss of the model is : 5.2119

Train log loss of the model is : 0.2579

Test f1 score of the model is : 79.3976

Train f1 score of the model is : 84.4911



- 1158 data points which belongs to fake news are classified correctly.
- 86 data points which belongs to real news category are misclassified as fake news.
- 899 data points which belongs to fake news category are misclassified as real news.
- 1898 data points which belongs to real news are classified correctly.

In [ ]:

## Decision Tree Classifier


```
In [90]: from sklearn.tree import DecisionTreeClassifier
#https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.Randomize
from sklearn.model_selection import RandomizedSearchCV
```

```
In [117... model = DecisionTreeClassifier()
```

```
In [118... hyperparameter = {"max_depth" : [3,5,7,9,11,13,15,17,19,21]}
search = RandomizedSearchCV(model, hyperparameter, scoring = "accuracy", random_stat
search.fit(x_tr,y_train)
```

Out[118...






```

In [119... plt.plot(hyperparameter["max_depth"], search.cv_results_['mean_test_score'])
plt.xlabel("Hyper parameter (max_depth)")
plt.ylabel("Accuracy score")
plt.title("Hyper parameter v/s Accuracy score")
plt.show()

```



```

In [120... search.cv_results_

```

```

Out[120... {'mean_fit_time': array([ 9.63867116, 11.72168837, 13.53531256, 14.75545163, 15.4415
7877,
        16.33877888, 17.65191264, 18.12520823, 18.52821202, 19.46276317])),
'std_fit_time': array([0.26887928, 0.19533027, 0.40742154, 0.26611401, 0.61157337,
        0.34445686, 0.55648399, 0.21417247, 0.41755096, 0.44823362])),
'mean_score_time': array([0.02299848, 0.02217112, 0.02457943, 0.01999717, 0.0206461
4,
        0.02136021, 0.0216804 , 0.02044678, 0.02290483, 0.02054877])),
'std_score_time': array([6.45042094e-03, 2.66942606e-03, 6.48007279e-03, 3.18948102
e-05,
        9.12627598e-04, 1.09920517e-03, 2.93394565e-03, 2.05235518e-03,
        1.80627660e-03, 2.77741039e-03])),
'param_max_depth': masked_array(data=[3, 5, 7, 9, 11, 13, 15, 17, 19, 21],
        mask=[False, False, False, False, False, False, False, False,
        False, False],
        fill_value='?',
        dtype=object),
'params': [{ 'max_depth': 3},
        { 'max_depth': 5},
        { 'max_depth': 7},
        { 'max_depth': 9},
        { 'max_depth': 11},
        { 'max_depth': 13},
        { 'max_depth': 15},
        { 'max_depth': 17},
        { 'max_depth': 19},
        { 'max_depth': 21}],
'split0_test_score': array([0.93257037, 0.94772657, 0.95515002, 0.96102691, 0.96381
07 ,
        0.96628518, 0.97030622, 0.97401794, 0.9758738 , 0.97618311])),
'split1_test_score': array([0.94370554, 0.96040829, 0.96628518, 0.97061553, 0.97525
518,

```

```

0.97834828, 0.98020414, 0.98082277, 0.98360656, 0.98453449]),
'split2_test_score': array([0.93595297, 0.95451733, 0.96008663, 0.96627475, 0.97029
703,
0.97277228, 0.97462871, 0.97524752, 0.97617574, 0.98019802]),
'split3_test_score': array([0.93997525, 0.95420792, 0.96039604, 0.96627475, 0.97029
703,
0.9737005 , 0.97710396, 0.97834158, 0.9789604 , 0.98112624]),
'split4_test_score': array([0.94059406, 0.95358911, 0.96163366, 0.96596535, 0.97122
525,
0.97339109, 0.97555693, 0.97834158, 0.97957921, 0.98112624]),
'mean_test_score': array([0.93855964, 0.95408984, 0.96071031, 0.96603146, 0.9701770
4,
0.97289947, 0.97555999, 0.97735428, 0.97883914, 0.98063362]),
'std_test_score': array([0.00388044, 0.00402148, 0.00355911, 0.00303914, 0.0036727
5,
0.00385585, 0.00323794, 0.00243103, 0.0027992 , 0.00267204]),
'rank_test_score': array([10, 9, 8, 7, 6, 5, 4, 3, 2, 1])}

```

In [121... search.best\_params\_

Out[121... {'max\_depth': 21}

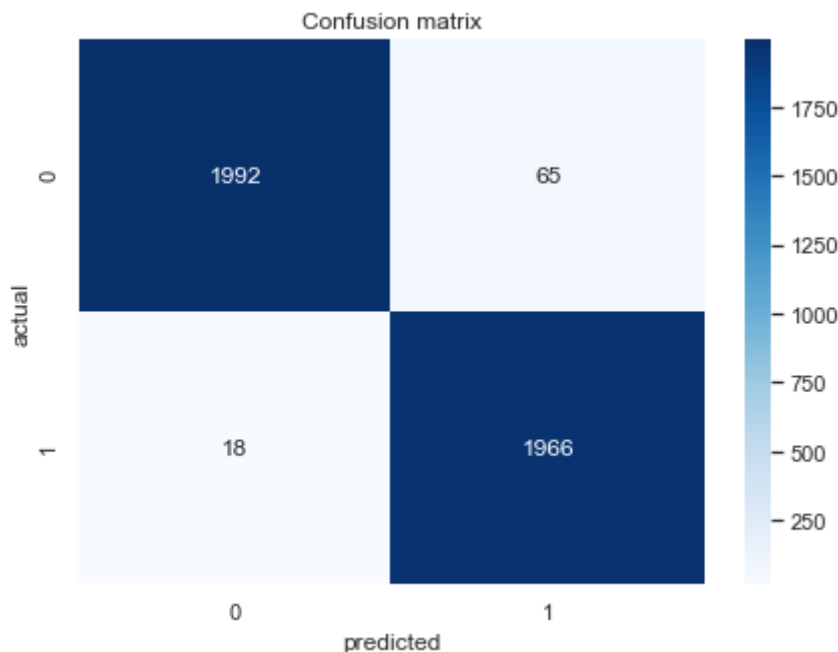
In [122... model = DecisionTreeClassifier(max\_depth = search.best\_params\_["max\_depth"])  
model.fit(x\_tr,y\_train)

Out[122... ▾ DecisionTreeClassifier  
DecisionTreeClassifier(max\_depth=21)

In [123... from sklearn.metrics import accuracy\_score, confusion\_matrix

In [124... y\_test\_pred = model.predict(x\_te)  
accuracy = accuracy\_score(y\_test,y\_test\_pred)  
print("Test accuracy of the model is : ", np.round(accuracy\*100,4), "%")  
train\_pred = model.predict(x\_tr)  
train\_accuracy = accuracy\_score(y\_train,train\_pred)  
print("Train accuracy of the model is : ", np.round(train\_accuracy\*100,4), "%")  
#=====
cm1 = confusion\_matrix(y\_test,y\_test\_pred)
plt.figure(figsize=(7,5))
sns.heatmap(cm1, annot=True,fmt="d",cmap='Blues')
plt.xlabel("predicted")
plt.ylabel("actual")
plt.title("Confusion matrix")
plt.show()

Test accuracy of the model is : 97.9461 %  
Train accuracy of the model is : 98.744 %



In [125...

```

y_test_pred = model.predict(x_te)
print("TEST PREDICTION: \n",y_test_pred[:3])
#predicting probability
y_test_pred_proba = model.predict_proba(x_te)
print("TEST PREDICTION PROBABILITY: \n",y_test_pred_proba[:3])

train_pred = model.predict(x_tr)
#predicting probability
train_pred_proba = model.predict_proba(x_tr)

accuracy = accuracy_score(y_test,y_test_pred)
print("Test log loss of the model is : ", np.round(accuracy*100,4), "%")
train_accuracy = accuracy_score(y_train,train_pred)
print("Train log loss of the model is : ", np.round(train_accuracy*100,4), "%")

print("=*50)

test_logloss = log_loss(y_test,y_test_pred_proba)
print("Test log loss of the model is : ", np.round(test_logloss,4))
train_logloss = log_loss(y_train,train_pred_proba)
print("Train log loss of the model is : ", np.round(train_logloss,4))

print("=*50)

test_f1 = f1_score(y_test,y_test_pred)
print("Test f1 score of the model is : ", np.round(test_f1*100,4))
train_f1 = f1_score(y_train,train_pred)
print("Train f1 score of the model is : ", np.round(train_f1*100,4))

cm1 = confusion_matrix(y_test,y_test_pred)
plt.figure(figsize=(7,5))
sns.heatmap(cm1, annot=True,fmt="d",cmap='Blues')
plt.xlabel("predicted")
plt.ylabel("actual")
plt.title("Confusion matrix")
plt.show()

```

TEST PREDICTION:


[1 1 1]

TEST PREDICTION PROBABILITY:

```

[[0.02566047 0.97433953]
 [0.02566047 0.97433953]
 [0.02566047 0.97433953]]
Test log loss of the model is : 97.9461 %
Train log loss of the model is : 98.744 %
=====
Test log loss of the model is : 0.3016
Train log loss of the model is : 0.0584
=====
Test f1 score of the model is : 97.9328
Train f1 score of the model is : 98.7206

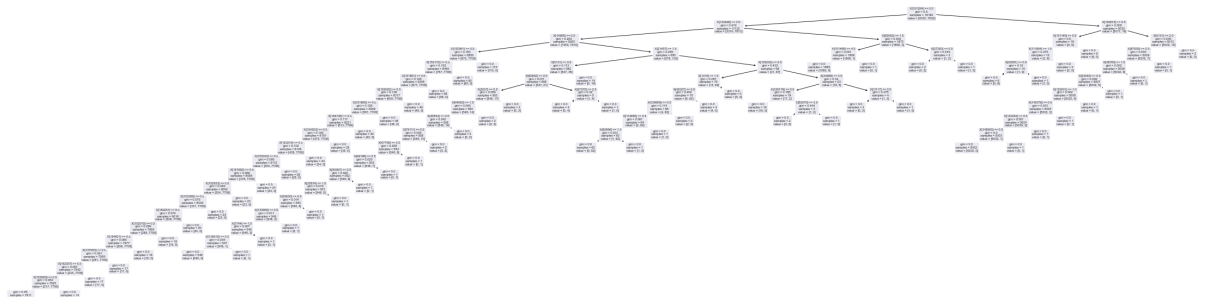
```



- 1996 data points which belongs to fake news are classified correctly.
- 17 data points which belongs to real news category are misclassified as fake news.
- 61 data points which belongs to fake news category are misclassified as real news.
- 1967 data points which belongs to real news are classified correctly.

```
In [126... from sklearn import tree
```

```
In [127... plt.figure(figsize=(40,10))
tree.plot_tree(model)
plt.show()
```



```
In [ ]:
```

```
In [169... model.feature_importances_
```


```
Out[169... array([0., 0., 0., ..., 0., 0., 0.])
```

```
In [170... coefs_df = pd.DataFrame()
coefs_df["coefs"] = model.feature_importances_
```

```

coefs_df["abs_coefs"] = abs(model.feature_importances_)
coefs_df.sort_values(by="abs_coefs", inplace=True, ascending=False)
top_coefs = coefs_df.head(50)
top_coefs
top_coefs["feature_index"] = top_coefs.index
top_coefs[["coefs"]].plot(kind="bar", figsize=(15, 6))
plt.title("Feature importance plot")
plt.xlabel("Feature index")
plt.ylabel("Co-efficient value")
plt.show()

```



## Random Forest Classifier

```

In [171...] from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import RandomizedSearchCV

```

```

In [172...] model = RandomForestClassifier()

```

```

In [173...] hyperparameter = {"n_estimators" : [50,100,150,200]}
search = RandomizedSearchCV(model, hyperparameter, scoring = "accuracy", random_stat
search.fit(x_tr,y_train)

```

```

Out[173...] RandomizedSearchCV
  estimator: RandomForestClassifier
    RandomForestClassifier

```

```

In [174...] search.cv_results_

```

```

Out[174...] {'mean_fit_time': array([132.15902462, 269.25143547, 401.6874825 , 491.48179626]),
'std_fit_time': array([ 1.58724338,  3.5085762 ,  4.6406422 , 54.04581438]),
'mean_score_time': array([0.70720143, 1.4028758 , 2.04618077, 2.38511906]),
'std_score_time': array([0.02881821, 0.03492622, 0.074619 , 0.61941251]),
'param_n_estimators': masked_array(data=[50, 100, 150, 200],
mask=[False, False, False, False],
fill_value='?',
dtype=object),
'params': [{'n_estimators': 50},
{'n_estimators': 100},
{'n_estimators': 150},
{'n_estimators': 200}],
'split0_test_score': array([0.9396845 , 0.94834519, 0.95112898, 0.95607795]),
'split1_test_score': array([0.94215899, 0.94989174, 0.94896381, 0.95484071]),

```


```
'split2_test_score': array([0.95080446, 0.95420792, 0.95575495, 0.95730198]),
'split3_test_score': array([0.93997525, 0.94461634, 0.94647277, 0.95482673]),
'split4_test_score': array([0.93966584, 0.94647277, 0.95420792, 0.95420792]),
'mean_test_score': array([0.94245781, 0.94870679, 0.95130569, 0.95545106]),
'std_test_score': array([0.00427563, 0.0032717 , 0.0033791 , 0.00110694]),
'rank_test_score': array([4, 3, 2, 1])}
```

In [175... search.best\_params\_

Out[175... {'n\_estimators': 200}

In [ ]: search.best\_estimator\_

```
In [176... plt.plot(hyperparameter["n_estimators"], search.cv_results_['mean_test_score'])
plt.xscale("log")
plt.xlabel("Hyper parameter (n_neighbors)")
plt.ylabel("Accuracy score")
plt.title("Hyper parameter v/s Accuracy score")
plt.show()
```



```
In [177... model = RandomForestClassifier(n_estimators = search.best_params_["n_estimators"])
model.fit(x_tr,y_train)
```

Out[177... **RandomForestClassifier**  
RandomForestClassifier(n\_estimators=200)

```
In [ ]: y_test_pred = model.predict(x_te)
accuracy = accuracy_score(y_test,y_test_pred)
print("Test accuracy of the model is : ", np.round(accuracy*100,4), "%")
train_pred = model.predict(x_tr)
train_accuracy = accuracy_score(y_train,train_pred)
print("Train accuracy of the model is : ", np.round(train_accuracy*100,4), "%")
#=====
cm2 = confusion_matrix(y_test,y_test_pred)
plt.figure(figsize=(7,5))
sns.heatmap(cm2, annot=True,fmt="d", cmap='Blues')
plt.xlabel("predicted")
plt.ylabel("actual")
plt.title("Confusion matrix")
plt.show()
```

- 2009 data points which belongs to fake news are classified correctly.

- 134 data points which belongs to real news category are misclassified as fake news.
- 48 data points which belongs to fake news category are misclassified as real news.
- 1850 data points which belongs to real news are classified correctly.

In [178...

```

y_test_pred = model.predict(x_te)
print("TEST PREDICTION: \n",y_test_pred[:3])
#predicting probability
y_test_pred_proba = model.predict_proba(x_te)
print("TEST PREDICTION PROBABILITY: \n",y_test_pred_proba[:3])

train_pred = model.predict(x_tr)
#predicting probability
train_pred_proba = model.predict_proba(x_tr)

accuracy = accuracy_score(y_test,y_test_pred)
print("Test log loss of the model is : ", np.round(accuracy*100,4), "%")
train_accuracy = accuracy_score(y_train,train_pred)
print("Train log loss of the model is : ", np.round(train_accuracy*100,4), "%")

print("="*50)

test_logloss = log_loss(y_test,y_test_pred_proba)
print("Test log loss of the model is : ", np.round(test_logloss,4))
train_logloss = log_loss(y_train,train_pred_proba)
print("Train log loss of the model is : ", np.round(train_logloss,4))

print("="*50)

test_f1 = f1_score(y_test,y_test_pred)
print("Test f1 score of the model is : ", np.round(test_f1*100,4))
train_f1 = f1_score(y_train,train_pred)
print("Train f1 score of the model is : ", np.round(train_f1*100,4))

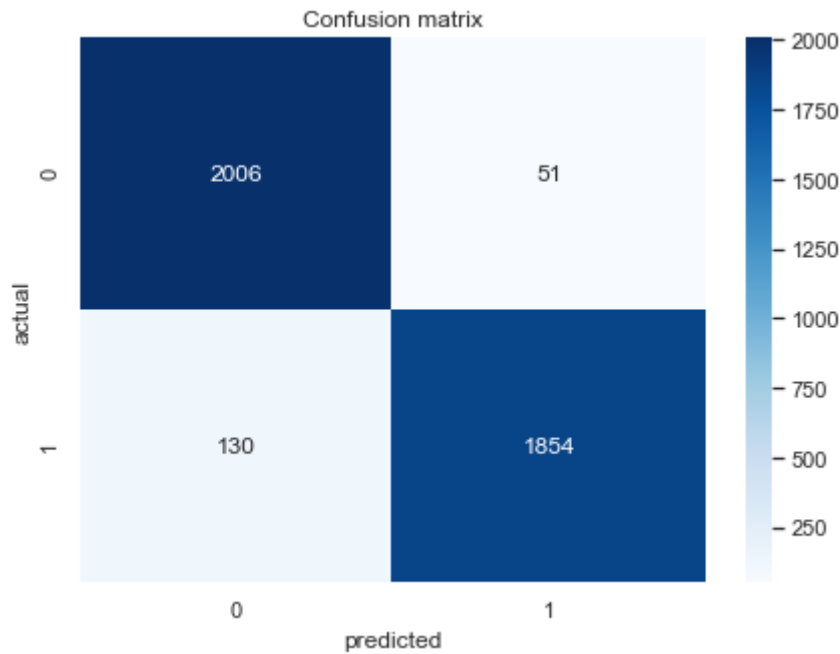
cm1 = confusion_matrix(y_test,y_test_pred)
plt.figure(figsize=(7,5))
sns.heatmap(cm1, annot=True,fmt="d", cmap='Blues')
plt.xlabel("predicted")
plt.ylabel("actual")
plt.title("Confusion matrix")
plt.show()

```

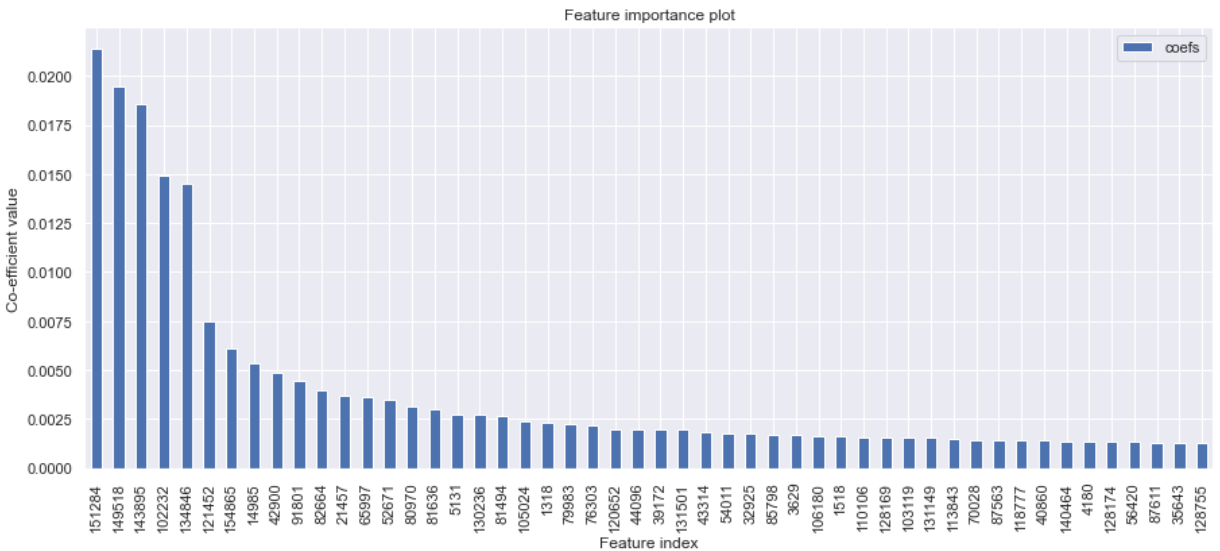
```

TEST PREDICTION:
[0 1 1]
TEST PREDICTION PROBABILITY:
[[0.585 0.415]
 [0.31  0.69 ]
 [0.165 0.835]]
Test log loss of the model is : 95.5209 %
Train log loss of the model is : 100.0 %
=====
Test log loss of the model is : 0.3061
Train log loss of the model is : 0.0974
=====
Test log loss of the model is : 95.3458
Train log loss of the model is : 100.0

```



```
In [179...
coefs_df = pd.DataFrame()
coefs_df["coefs"] = model.feature_importances_
coefs_df["abs_coefs"] = abs(model.feature_importances_)
coefs_df.sort_values(by="abs_coefs", inplace=True, ascending=False)
top_coefs = coefs_df.head(50)
top_coefs
top_coefs["feature index"] = top_coefs.index
top_coefs[["coefs"]].plot(kind="bar",figsize=(15, 6))
plt.title("Feature importance plot")
plt.xlabel("Feature index")
plt.ylabel("Co-efficient value")
plt.show()
```



```
In [187...
In [ ]:
In [ ]:
In [ ]:
```